

1. What did you find easy to do in Dash?

Creating interactive visualizations and layouts in Dash was straightforward due to its intuitive component-based structure and seamless integration with Plotly. The `dcc` (Dash Core Components) and `html` modules made it easy to build a tabbed interface with dropdowns, inputs, buttons, and graphs, as seen in the Iris classifier app. Defining callbacks to handle user interactions, such as updating scatter plots or histograms based on dropdown selections, was particularly simple because Dash's callback decorator clearly maps inputs to outputs. The ability to leverage Plotly Express for quick, customizable visualizations (e.g., `px.scatter` and `px.histogram`) further simplified the process of displaying the Iris dataset's features dynamically.

2. What was hard to implement or you didn't wind up getting it to work?

Handling backend API integration and error responses proved challenging, particularly ensuring robust communication between the Dash front-end and the Flask backend. In the provided app, issues like the "Invalid dataset index" and "history" errors highlighted difficulties in aligning the front-end callbacks (e.g., `update_output_train` and `update_output_score`) with the backend's response formats. Debugging these required extensive validation of inputs and response structures, which was complicated by the lack of visibility into the backend's `retrain_model` function. Implementing graceful error handling for non-JSON responses or unexpected backend behavior was also tricky, as it required anticipating various failure modes without modifying the backend.

3. What other components, or what revised dashboard design, would you suggest to better assist in explaining the behavior of the Iris model to a client?

To better explain the Iris model's behavior, I'd suggest adding a **Model Performance** tab with components like a confusion matrix (using `dcc.Graph` with `plotly.figure_factory.create_annotated_heatmap`), ROC curves, and precision-recall curves to visualize classification performance. A `dcc.Store` component could cache model metrics and predictions for cross-tab access. Incorporating a `dash_table.DataTable` to display feature importance or SHAP values would clarify which features (e.g., petal length) drive predictions. Design-wise, a sidebar layout with collapsible sections for data exploration, training, and evaluation could improve navigation, and tooltips (via

`dcc.Tooltip`) could provide client-friendly explanations of technical terms like “dataset ID” or “model retraining.”

4. Can you think of better ways to link the “back end” Iris model and its results with the front-end Dash functions?

A more robust approach to linking the backend Iris model with the Dash front-end would involve using WebSockets (via `dash_extensions` or `Flask-SocketIO`) for real-time updates, enabling asynchronous training progress feedback or live scoring results. Implementing a REST API client wrapper in Python (e.g., using `requests.Session`) within the Dash app could centralize API calls, handle authentication, and standardize error handling. Additionally, defining a clear API contract (e.g., with OpenAPI/Swagger) would ensure consistent response formats (e.g., always returning JSON with predictable keys like “history” or “prediction”). Using `dcc.Store` to cache backend responses locally could reduce redundant API calls and improve performance, especially for frequently accessed model metrics or dataset IDs.